

Roll no:COBA020

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import datetime as dt
```

In [2]:

```
df = pd.read_csv("uber.csv")
df.head()
```

Out[2]:

Unnamed: key fare\_amount pickup\_datetime pickup\_longitude pickup\_latitude

			2015-05-07		2015-05-07
0	24238194	7.5	-73.999817	40.7	
			19:52:06.0000003		19:52:06 UTC
			2009-07-17		2009-07-17
1	27835199	7.7	-73.994355	40.7	
			20:04:56.0000002		20:04:56 UTC
			2009-08-24		2009-08-24
2	44984355	12.9	-74.005043	40.74	
			21:45:00.00000061		21:45:00 UTC
			2009-06-26		2009-06-26
3	25894730	5.3	-73.976124	40.7	
			08:22:21.0000001		08:22:21 UTC
			2014-08-28		2014-08-28
4	17610152	16.0	-73.925023	40.74	
			17:47:00.000000188		17:47:00 UTC

In [3]:

```
df.drop(columns=['Unnamed: 0', 'key'], inplace=True)
```

In [4]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
#   ...
```

```

---  -----  -----  0
fare_amount      200000 non-null  float64
1  pickup_datetime  200000 non-null  object
2  pickup_longitude  200000 non-null  float64
3  pickup_latitude   200000 non-null  float64
4  dropoff_longitude  199999 non-null  float64
5  dropoff_latitude   199999 non-null  float64
6  passenger_count    200000 non-null  int64  dtypes: float64(5), int64(1),
object(1) memory usage: 10.7+ MB

```

## Dropping null rows

In [5]:

```
df.dropna(how='any',inplace=True)
```

In [6]:

```
df.isnull().sum()
```

Out[6]:

```

fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude   0
dropoff_longitude 0
dropoff_latitude  0
passenger_count  0
dtype: int64




```




Roll no:COBA020

## Boxplots

In [7]:

```
for col in df.select_dtypes(exclude=['object']):  
    plt.figure()  
    sns.boxplot(data=df,x=col)
```





## Dropping outliers

**-90 < latitude < 90 -**

**180 < longitude < 180**

**fare > 0**

**0 < passenger\_count < 50** In [8]:

```
df = df[
```

Roll no:COBA020

```
(df.pickup_latitude > -90) & (df.pickup_latitude < 90) &
(df.dropoff_latitude > -90) & (df.dropoff_latitude < 90) &
(df.pickup_longitude > -180) & (df.pickup_longitude < 180) &
(df.dropoff_longitude > -180) & (df.dropoff_longitude < 180) &
(df.fare_amount > 0) & (df.passenger_count > 0) & (df.passenger_count < 50)
]
```

## Calculating Distance

In [9]:

```
from math import cos, asin, sqrt, pi
import numpy as np

def distance(lat_1,lon_1,lat_2,lon_2):
#     lat1 = row.pickup_latitude
#     lon1 = row.pickup_longitude
#     lat2 = row.dropoff_latitude
#     lon2 = row.dropoff_longitude
    lon_1, lon_2, lat_1, lat_2 = map(np.radians, [lon_1, lon_2, lat_1, lat_2]) #Degrees

    diff_lon = lon_2 - lon_1
    diff_lat = lat_2 - lat_1

    km = 2 * 6371 * np.arcsin(np.sqrt(np.sin(diff_lat/2.0)**2 + np.cos(lat_1) * np.cos(

    return km
```

In [10]:

```
temp = distance(df['pickup_latitude'],df['pickup_longitude'],df['dropoff_latitude'],df['
temp.head()
```

Out[10]:

```
0    1.683323
1    2.457590
2    5.036377
3    1.661683 4    4.475450 dtype: float64
```

In [11]:

```
df_new = df.copy()
df_new['Distance'] = temp
df = df_new
df.head()
```

Out[11]:


	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropo
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	

In [12]:

```
sns.boxplot(data=df,x='Distance')
```

Out[12]: <AxesSubplot:

xlabel='Distance'>



In [13]:

```
df = df[(df['Distance'] < 200) & (df['Distance'] > 0)]
```

## Date and Time features extract

In [14]:

```
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
```

<ipython-input-14-834f97bbe4ec>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
```

In [15]:

```
df['week_day'] = df['pickup_datetime'].dt.day_name()
df['Year'] = df['pickup_datetime'].dt.year
df['Month'] = df['pickup_datetime'].dt.month
df['Hour'] = df['pickup_datetime'].dt.hour
```

<ipython-input-15-b91c1da9c026>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)) df['week\_day'] =

```
df['pickup_datetime'].dt.day_name()
```

<ipython-input-15-b91c1da9c026>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)) df['Year'] = df['pickup\_datetime'].dt.year

<ipython-input-15-b91c1da9c026>:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)) df['Month'] = df['pickup\_datetime'].dt.month

<ipython-input-15-b91c1da9c026>:4: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)) df['Hour'] = df['pickup\_datetime'].dt.hour

In [16]:

```
df.drop(columns=['pickup_datetime','pickup_latitude','pickup_longitude','dropoff_latitude',
```

<ipython-input-16-a7c1789815f4>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.drop(columns=['pickup_datetime','pickup_latitude','pickup_longitude',  
'dropoff_latitude','dropoff_longitude'],inplace=True)
```



In [17]:

```
df.head()
```

Out[17]:

	fare_amount	passenger_count	Distance	week_day	Year	Month	Hour
0	7.5	1	1.683323	Thursday	2015	5	
	19						
1	7.7	1	2.457590	Friday	2009	7	20
2	12.9	1	5.036377	Monday	2009	8	21
3	5.3	3	1.661683	Friday	2009	6	8
4	16.0	5	4.475450	Thursday	2014	8	17

In  
[18]:

```
temp = df.copy()

def convert_week_day(day):
    if day in ['Monday', 'Tuesday', 'Wednesday', 'Thursday']:
        return 0 # Weekday
    return 1 # Weekend

def convert_hour(hour):
    if 5 <= hour <= 12:
        return 1
    elif 12 < hour <= 17:
        return 2
    elif 17 < hour < 24:
        return 3
    return 0

df['week_day'] = temp['week_day'].apply(convert_week_day)
df['Hour'] = temp['Hour'].apply(convert_hour)
df.head()
```

<ipython-input-18-655f90749f34>:17: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)) df['week\_day'] = temp['week\_day'].apply(convert\_week\_day)

<ipython-input-18-655f90749f34>:18: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)) df['Hour'] = temp['Hour'].apply(convert\_hour)

Out[18]:

	fare_amount	passenger_count	Distance	week_day	Year	Month	Hour
0	7.5	1	1.683323	0	2015	5	3
1	7.7	1	2.457590	1	2009	7	3
	12.9	1	5.036377	0	2009	8	3
3	5.3	3	1.661683	1	2009	6	1

4 16.0 5 4.475450 0 2014 8 2

## Correlation Matrix

In [19]:

```
df.corr()
```


Out[19]:

	fare_amount	passenger_count	Distance	week_day	Year	Month
fare_amount	1.000000	0.011884	0.778667	0.002305	0.120430	0.024120
passenger_count	0.011884	1.000000	0.005112	0.035882	0.005339	0.008818
Distance	0.778667	0.005112	1.000000	0.014518	0.018617	0.007373
week_day	0.002305	0.035882	0.014518	1.000000	0.006910	-0.007328
Year	0.120430	0.005339	0.018617	0.006910	1.000000	-0.115182
Month	0.024120	0.008818	0.007373	-0.007328	-0.115182	1.000000
Hour	-0.021078	0.013572	-0.022691	-0.078129	0.001131	-0.005410

In [20]:

```
sns.scatterplot(y=df['fare_amount'],x=df['Distance'])
```

Out[20]: <AxesSubplot: xlabel='Distance',  
ylabel='fare\_amount'>



In [21]:

```
from sklearn.preprocessing import StandardScaler  
x = df[['Distance']].values  
y = df['fare_amount'].values.reshape(-1,1)
```

In

**Independent Variable: Distance**

**Dependent Variable: fare\_amount**

[31]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state=10) In [32]:
```

```
std_x = StandardScaler()
x_train = std_x.fit_transform(x_train)
```

In [33]:

```
x_test = std_x.transform(x_test)
```

In [34]:

```
std_y = StandardScaler()
y_train = std_y.fit_transform(y_train)
```

In [35]:

```
y_test = std_y.transform(y_test)
```

In [36]:

```
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
def fit_predict(model):
    model.fit(x_train, y_train.ravel())
    y_pred = model.predict(x_test)
    r_squared = r2_score(y_test, y_pred)
    RMSE = mean_squared_error(y_test, y_pred, squared=False)
    MAE = mean_absolute_error(y_test, y_pred)
    print('R-squared: ', r_squared)
    print('RMSE: ', RMSE)
    print("MAE: ", MAE)
```

In [37]:

```
from sklearn.linear_model import LinearRegression
```

In [38]:

```
fit_predict(LinearRegression())
```

R-squared: 0.604116792084117

RMSE: 0.6290054895695945

MAE: 0.27552329590959823 In

[39]:

```
from sklearn.ensemble import RandomForestRegressor  
fit_predict(RandomForestRegressor())
```

R-squared: 0.652350257870196

RMSE: 0.589443049630681

MAE: 0.2921068537600526

In [ ]: